

Bloom Filters & Count-Min Sketch

Anil Maheshwari

School of Computer Science
Carleton University
Canada

August 6, 2018

- 1 Bloom Filter
- 2 An Interview Problem
- 3 Count-Min Sketch

Problem Definition

Let U be the universe.

Input: A subset $S \subseteq U$.

Query: For any $q \in U$, decide whether $q \in S$.

Objective

Answer queries quickly and use very little extra space.

SPAM Detection

U = All possible email addresses;

S = My collection of non-junk email addresses.

Query: Given any $q \in U$, report whether $q \in S$.

- Proposed by Bloom in CACM 1970 - *Space/Time tradeoffs in Hash Coding with Allowable Errors*. (7000 Citations)
- Space-Efficient Probabilistic Data Structure for Membership Testing
- May have false positives
- Numerous Variants: Counting Filters, Dynamic Filters with insertion/deletion of elements in S ,
- Vast Applications: Estimating size of union/intersection of sets, Avoid caching 'one-hit wonders', Google Bigtable, Chrome's uses it to detect malicious URLs.
- Refined Analysis in 2008 by Bose et al.

Data Structure

An array B consisting of m bits and k hash functions h_1, h_2, \dots, h_k , where $h_i : U \rightarrow \{1, \dots, m\}$

Initialization

$B \leftarrow 0$.

For all $x \in S$, set

$B[h_1(x)] = B[h_2(x)] = \dots = B[h_k(x)] = 1$.

Answering Query

For any query $q \in U$,
if $B[h_1(q)] = B[h_2(q)] = \dots = B[h_k(q)] = 1$, report $q \in S$,
else report $q \notin S$.

Observation

If $q \in S$, the queries are answered correctly.

False Positives

Suppose $q \notin S$

If $B[h_1(q)] = B[h_2(q)] = \dots = B[h_k(q)] = 1$,
we will report that $q \in S$.

Estimating Probability of False-Positives

Assume $n = |S|$.

- nk times, we attempt to set locations in B to "1".
- What is the probability that $B[l] = 1$?
- Complementary Event: $Pr(B[l] = 0) = (1 - \frac{1}{m})^{nk}$
- $p = Pr(B[l] = 1) = 1 - (1 - \frac{1}{m})^{nk}$
- For False-Positive to occur, all of the k specified locations $B[h_1(q)], \dots, B[h_k(q)]$ must be "1".

Bloom70

$$Pr(B[h_1(q)] = B[h_2(q)] = \dots = B[h_k(q)] = 1) = p^k.$$

Estimating Probability of False-Positives

Assume $n = |S|$.

- nk times, we attempt to set locations in B to "1".
- What is the probability that $B[l] = 1$?
- Complementary Event: $Pr(B[l] = 0) = (1 - \frac{1}{m})^{nk}$
- $p = Pr(B[l] = 1) = 1 - (1 - \frac{1}{m})^{nk}$
- For False-Positive to occur, all of the k specified locations $B[h_1(q)], \dots, B[h_k(q)]$ must be "1".

Bloom70

$$Pr(B[h_1(q)] = B[h_2(q)] = \dots = B[h_k(q)] = 1) = p^k.$$

Estimating Probability of False-Positives

Assume $n = |S|$.

- nk times, we attempt to set locations in B to "1".
- What is the probability that $B[l] = 1$?
- Complementary Event: $Pr(B[l] = 0) = (1 - \frac{1}{m})^{nk}$
- $p = Pr(B[l] = 1) = 1 - (1 - \frac{1}{m})^{nk}$
- For False-Positive to occur, all of the k specified locations $B[h_1(q)], \dots, B[h_k(q)]$ must be "1".

Bloom70

$$Pr(B[h_1(q)] = B[h_2(q)] = \dots = B[h_k(q)] = 1) = p^k.$$

Estimating Probability of False-Positives

Assume $n = |S|$.

- nk times, we attempt to set locations in B to "1".
- What is the probability that $B[l] = 1$?
- Complementary Event: $Pr(B[l] = 0) = (1 - \frac{1}{m})^{nk}$
- $p = Pr(B[l] = 1) = 1 - (1 - \frac{1}{m})^{nk}$
- For False-Positive to occur, all of the k specified locations $B[h_1(q)], \dots, B[h_k(q)]$ must be "1".

Bloom70

$$Pr(B[h_1(q)] = B[h_2(q)] = \dots = B[h_k(q)] = 1) = p^k.$$

Let $n = 1$, $m = 2$, $k = 2$,
 $U = \{x, y\}$, $S = \{x\}$ and $q = y \neq x$.

After Initialization B has the following configuration:

B	Pr. of specific config. of B	Given B , Cond. Pr. that $B[h_1(y)] = B[h_2(y)] = 1$		
<table border="1"><tr><td>1</td><td>0</td></tr></table>	1	0	$1/2 \times 1/2 = 1/4$	$1/2 \times 1/2 = 1/4$
1	0			
<table border="1"><tr><td>0</td><td>1</td></tr></table>	0	1	$1/2 \times 1/2 = 1/4$	$1/2 \times 1/2 = 1/4$
0	1			
<table border="1"><tr><td>1</td><td>1</td></tr></table>	1	1	$2 \times 1/2 \times 1/2 = 1/2$	$1 \times 1 = 1$
1	1			

Since the three rows are mutually exclusive, the probability of False-Positive is

$$1/4 \times 1/4 + 1/4 \times 1/4 + 1/2 \times 1 = 10/16.$$

Bloom Filter

An Interview
Problem

Count-Min Sketch

An Example Contd.

$$n = 1, m = 2, k = 2.$$

Note that Bloom's result states that the probability of false-positive is p^k , where $p = 1 - (1 - \frac{1}{m})^{kn}$.

From Bloom's computation,

$$p = 1 - (1 - \frac{1}{m})^{kn} = 1 - (1 - \frac{1}{2})^{2 \times 1} = 3/4, \text{ and}$$
$$p^k = p^2 = 9/16.$$

But,

$$9/16 \neq 10/16$$

An Example Contd.

$$n = 1, m = 2, k = 2.$$

Note that Bloom's result states that the probability of false-positive is p^k , where $p = 1 - (1 - \frac{1}{m})^{kn}$.

From Bloom's computation,

$$p = 1 - (1 - \frac{1}{m})^{kn} = 1 - (1 - \frac{1}{2})^{2 \times 1} = 3/4, \text{ and}$$
$$p^k = p^2 = 9/16.$$

But,

$$9/16 \neq 10/16$$

We came up with a fairly technical proof and showed that

Theorem

Let $p_{k,n,m}$ be the false-positive rate for a Bloom filter that stores n elements of a set S in a bit-vector of size m using k hash functions.

- 1 We can express $p_{k,n,m}$ in terms of the Stirling number of second kind as follows:

$$p_{k,n,m} = \frac{1}{m^{k(n+1)}} \sum_{i=1}^m i^k i! \binom{m}{i} \left\{ \begin{matrix} kn \\ i \end{matrix} \right\}$$

- 2 Let $p = 1 - (1 - 1/m)^{kn}$, $k \geq 2$ and $\frac{k}{p} \sqrt{\frac{\ln m - 2k \ln p}{m}} \leq c$ for some $c < 1$. Upper and lower bounds on $p_{k,n,m}$ are given by

$$p^k < p_{k,n,m} \leq p^k \left(1 + O\left(\frac{k}{p} \sqrt{\frac{\ln m - 2k \ln p}{m}} \right) \right)$$

Summary

- 1 A simple scheme for testing membership.
Has one-sided error, i.e., false positives.
- 2 How to find the right number of hash functions and right size of the filter?
- 3 Implemented in various search engines, routers, SPAM filters, . . .
- 4 Unpleasant analysis.
- 5 Challenge: A nicer analysis.
Hopefully, this will help with the analysis of variants of Bloom Filters.

Finding the Majority Element

Input: A stream consisting of n elements and it is given that it has a majority element.

Output: The majority element.

- Store the stream in an array A . Sort and pick the middle element (if elements can be ordered).
- Count frequency of each element.
- What if we can only use $O(1)$ space!

Finding the Majority Element

Input: A stream consisting of n elements and it is given that it has a majority element.

Output: The majority element.

- Store the stream in an array A . Sort and pick the middle element (if elements can be ordered).
- Count frequency of each element.
- What if we can only use $O(1)$ space!

Finding the Majority Element

Input: A stream consisting of n elements and it is given that it has a majority element.

Output: The majority element.

- Store the stream in an array A . Sort and pick the middle element (if elements can be ordered).
- Count frequency of each element.
- What if we can only use $O(1)$ space!

Finding the Majority Element

Input: A stream consisting of n elements and it is given that it has a majority element.

Output: The majority element.

- Store the stream in an array A . Sort and pick the middle element (if elements can be ordered).
- Count frequency of each element.
- What if we can only use $O(1)$ space!

Majority Algorithm

Input: Array A of size n consisting a majority element

Output: The majority element

```
1   $c \leftarrow 0$ 
2  for  $i = 1$  to  $n$  do
3      if  $c = 0$  then
4           $current \leftarrow A[i]; c \leftarrow c + 1$ 
5      end
6      else
7          if  $A[i] = current$  then
8               $c \leftarrow c + 1$ 
9          end
10         else
11              $c \leftarrow c - 1$ 
12         end
13     end
14 end
15 return  $current$ 
```

Observations

- 1 Algorithm maintains only two variables: c and current.
- 2 Correctness: Each non-majority element can 'kill' at most one majority element.

Generalize

For a data stream, using very little space, we are interested to report

- 1 All the elements that occur frequently, e.g. say at least 2% of times.
- 2 For each element, its (approximate) frequency.

Count-Min Sketch Data Structure

Input: An array A consisting of n natural numbers and r hash functions h_1, \dots, h_r , where $h_i : \mathbb{N} \rightarrow \{1, \dots, b\}$

Output: $CMS[\cdot, \cdot]$ table consisting of r rows and b columns

```
1 for  $i = 1$  to  $r$  do
2   |   for  $j = 1$  to  $b$  do
3     |    $CMS[i, j] \leftarrow 0$ 
4   |   end
5 end
6 for  $i = 1$  to  $n$  do
7   |   for  $j = 1$  to  $r$  do
8     |    $CMS[j, h_j(A[i])] \leftarrow CMS[j, h_j(A[i])] + 1$ 
9   |   end
10 end
11 return  $CMS[\cdot, \cdot]$ 
```

Updating CMS table

An example with $b = 10$ and $r = 3$ and assume that stream $A = xyy$

After Initialization:

	1	2	3	4	5	6	7	8	9	10
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0

Updating CMS table

Insertion of x : $h_1(x) = 3$, $h_2(x) = 8$, and $h_3(x) = 5$:

	1	2	3	4	5	6	7	8	9	10
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0

After inserting x :

	1	2	3	4	5	6	7	8	9	10
1	0	0	1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	1	0	0
3	0	0	0	0	1	0	0	0	0	0

Updating CMS table

Insertion of 1st y : $h_1(y) = 6$, $h_2(y) = 8$, and $h_3(y) = 1$ that hashes to locations 6, 8, and 1:

	1	2	3	4	5	6	7	8	9	10
1	0	0	1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	1	0	0
3	0	0	0	0	1	0	0	0	0	0

After insertion of 1st y :

	1	2	3	4	5	6	7	8	9	10
1	0	0	1	0	0	1	0	0	0	0
2	0	0	0	0	0	0	0	2	0	0
3	1	0	0	0	1	0	0	0	0	0

Updating CMS table

Insertion of 2nd y (hashes to same locations 6,8, and 1):

	1	2	3	4	5	6	7	8	9	10
1	0	0	1	0	0	1	0	0	0	0
2	0	0	0	0	0	0	0	2	0	0
3	1	0	0	0	1	0	0	0	0	0

After inserting 2nd y :

	1	2	3	4	5	6	7	8	9	10
1	0	0	1	0	0	2	0	0	0	0
2	0	0	0	0	0	0	0	3	0	0
3	2	0	0	0	1	0	0	0	0	0

n = #items in the stream.

f_x^* = frequency of x in the stream.

Let $f_x = \min\{CMS[1, h_1(x)], \dots, CMS[r, h_r(x)]\}$

- 1 The size of CMS table ($= br$) is independent of n .
- 2 CMS table can be computed in $O(br + nr)$ time.
- 3 For any $x \in A$, and for any $j = 1, \dots, r$,
 $CMS[j, h_j(x)] \geq f_x^*$.
- 4 Therefore, $f_x \geq f_x^*$ (i.e., f_x is an overestimate).

Claim

Let $b = \frac{2}{\epsilon}$. Then $Pr[|f_x - f_x^*| \geq \epsilon n] \leq \frac{1}{2^r}$

Proof Sketch: Let V be the set of different values in the stream A . Define indicator r.v. I_y corresponding to each value $y \in A$ as follows:

$$I_y = \begin{cases} 1 & \text{if } h_j(y) = h_j(x) \\ 0, & \text{otherwise} \end{cases}$$

Note: $Pr(I_y = 1) = 1/b$, $E[I_y] = 1/b$.

Observe,

$$CMS[j, h_j(x)] = f_x^* + \sum_{\substack{y \in V \\ y \neq x}} I_y * f_y^* \quad (1)$$

Bloom Filter

An Interview
Problem

Count-Min Sketch

- By setting $b = \frac{2}{\epsilon}$, we obtain

$$E[CMS[j, h_j(x)]] \leq f_x^* + n/b = f_x^* + \epsilon n/2 \quad (2)$$

- Let $X_j = |CMS[j, h_j(x)] - f_x^*|$
- Then, $E[X_j] \leq n/b = \epsilon n/2$.
- Recall Markov's inequality - Pr. that a r.v. deviates from its expectation by a factor of c is at most $1/c$.
- Thus, $Pr(X_j > 2(\epsilon n/2)) \leq 1/2$.
- This also holds for each value of $j = 1, \dots, r$.
- Furthermore, X_j is independent of X_k as hash functions h_j and h_k are independent for any $k \neq j$.
- Therefore, for

$$f_x = \min\{CMS[1, h_1(x)], \dots, CMS[r, h_r(x)]\},$$
$$Pr[|f_x - f_x^*| \geq \epsilon n] \leq \frac{1}{2^r} \quad \square$$

Corollary

We have that $f_x^* \leq f_x \leq f_x^* + \epsilon n$ with probability at least $1 - 1/2^r$.

Suppose we are interested to report all elements in A that occur approx. n/k times for some integer k .

- Set $\epsilon = 1/3k$. Then, $b = 2/\epsilon = 6k$.
- Size of CMS table is $br = 6kr$.
- Scan A and update the *CMS* table as before.
- Also, maintain a set of $O(k)$ items that occur most frequently among all the elements in A scanned so far.
- The items are stored in a heap with their key as their f_x value.
- Assume we have scanned $i - 1$ items and have updated the *CMS* table and the heap.

Consider the i -th item (say $x = A[i]$) and we perform the following:

- ① Update the counts in the CMS table by executing $CMS[j, h_j(x)] \leftarrow CMS[j, h_j(x)] + 1$, for $j = 1$ to r .
- ② Let $f_x = \min\{CMS[1, h_1(x)], \dots, CMS[r, h_r(x)]\}$. If $f_x \geq i/k$, we perform the following heap operations:
 - ① If $x \in \text{heap}$, delete x and re-insert it again with the updated f_x value.
 - ② If $x \notin \text{heap}$, then we insert it in the heap, but remove all the elements whose count is less than i/k .

Claim

[Cormode and Muthukrishnan 2005] Elements that occur approx. n/k times in a data stream of size n can be reported in $O(kr + nr + n \log k)$ time using $O(kr)$ space with high probability.

Proof Sketch:

- Note: $f_x^* \leq f_x \leq f_x^* + \epsilon n = f_x^* + n/3k$.
- Thus, heap contains elements whose frequency is at least $n/k - n/3k = 0.667n/k$ with high probability.
- Size of heap = $O(k)$
- Total Time = $O(br + nr + n \log k) = O(kr + nr + n \log k)$, as $b = 6k$.
- Total Space = $O(br + k) = O(kr)$ \square

- Simple idea with important applications in

① Compressed Sensing, Anomaly Detection, Finding Outliers, ...

② Sketch Data Structure:

Consider a vector $v = (v_1, v_2, \dots, v_n)$.

Initially $v = 0$.

Update at time t is a pair (j, c) : $v_j \leftarrow v_j + c$.

Using only small space, answer queries of the form

- ① Point Query: Report v_i
- ② Range Query $[l, r]$: Report $\sum_{i=l}^r v_i$
- ③ Inner product of two vectors: $u \cdot v$
- ④ In general c can be positive or negative - replace min by median.

References

- 1 *Space/Time Trade-offs in Hash Coding with Allowable Errors*, B. H. Bloom, Communications of ACM, 13 (7): 422-426, 1970.
- 2 *An improved data stream summary: the count-min sketch and its applications*, G. Cormode and S. Muthukrishnan, J. Algorithms 55(1): 58-75, 2005.
- 3 *On the false-positive rate of Bloom filters*, Bose et al., Inf. Process. Letters 108(4): 210-213, 2008.
- 4 mmds.org - Mining Massive Data Sets by Leskovec, Rajaraman, and Ullman.

Thank-you