

STL Containers

Reading: Eckel, Vol. 2
Ch. 7 Generic Containers

The C++ Standard Template Library (STL)

The STL is a major component of the C++ Standard Library; it is a large collection of general-purpose generic classes, functions, and iterators:

1. Generic **containers** that take the element type as a parameter.
- e.g., `vector`, `list`, `deque`, `set`, `stack`, `queue`, ...
2. Different kinds of **iterators** that can navigate through the containers.
3. **Algorithms** that (via iterators) perform an interesting operation on a range of elements.
- e.g., `sort`, `random_shuffle`, `transform`, `find`



Design Philosophy of the STL

Generic **containers** that take the element type as a parameter.

- know (almost) nothing about the element type
exception: ordered containers expect elements to have `operator<`
- operations are (mostly) limited to add, remove, and retrieve
- define their own **iterators**

Useful, efficient, generic **algorithms** that:

- know nothing about the data structures they operate on
- know (almost) nothing about the elements in the structures
- operate on range of elements accessed via **iterators**



Design Philosophy of the STL

- STL algorithms are designed so that (almost) any algorithm can be used with any STL container, or any other data structure that supports iterators.
 - Element type must support copy constructor/assignment.
- For **ordered** containers, the element type must support `operator<` or you can provide a special *functor* (function-object) of your own.
- The STL assumes *value semantics* for its contained elements: elements are *copied* to/from containers more than you might realize.
- The container methods and algorithms are highly efficient; it is unlikely that you could do better.



No Inheritance in the STL!

Basically, the primary designer (Alexander Stepanov) thinks that OOP (i.e., inheritance) is wrong, and generic programming is better at supporting polymorphism, flexibility and reuse.

- Templates provide a **more flexible** (“ad hoc”) kind of polymorphism.
- The containers are different enough that code reuse isn't really practical.
- Container methods are **not virtual**, to improve efficiency.



STL References

There are good on-line references:

C++ Reference

<http://www.cplusplus.com/reference/stl/>
<http://www.cplusplus.com/reference/algorithm/>

SGI Standard Template Library Programmer's Guide

<http://www.sgi.com/tech/stl/>



Review: Polymorphic Containers

Suppose we want to model a graphical Scene that has an ordered list of Figures (i.e., Rectangles, Circles, and maybe other concrete classes we haven't implemented yet).

- `Figure` is an abstract base class (ABC)
- `Rectangle`, `Circle`, etc. are derived classes

What should the list look like?

1. `vector<Figure>`
2. `vector<Figure&>`
3. `vector<Figure*>`



Containers of Objects or Pointers?

```
Circle c ("red");  
vector<Figure> figList;  
figList.emplace_back(c);
```

Objects:

- copy operations could be expensive
- two red circles
- changes to one do not affect the other
- when `figList` dies, it will destroy its copy of red circle
- risk of static slicing

```
Circle c ("red");  
vector<Figure*> figList;  
figList.emplace_back(&c);
```

Pointers:

- allows for polymorphic containers
- when `figList` dies, only pointers are destroyed
- client code must cleanup referents of pointer elements



```
#include <iostream>
#include <string>
#include <vector>
using namespace std;
class Balloon {
public:
    Balloon (string colour);
    Balloon (const Balloon& b); // Copy constructor
    virtual ~Balloon();
    virtual void speak() const;
private:
    string colour;
};
Balloon::Balloon(string colour) : colour{colour} {
    cout << colour << " balloon is born" << endl;
}
Balloon::Balloon(const Balloon& b) : colour{b.colour} {
    cout << colour << " copy balloon is born" << endl;
}
void Balloon::speak() const {
    cout << "I am a " << colour << " balloon" << endl;
}
Balloon::~Balloon() {
    cout << colour << " balloon dies" << endl;
}
```



```
// How many Balloons are created?
int main (int argc, char* argv[]) {
    vector<Balloon> v;
    Balloon rb ("red");
    v.push_back(rb);
    Balloon gb ("green");
    v.push_back(gb);
    Balloon bb ("blue");
    v.push_back(bb);
}
```



STL Containers



STL Containers: Conceptual View

C++98/03 defines three main data-container categories:


1. Sequence containers: vector, deque, list
2. Container adapters: stack, queue, priority_queue
3. Ordered associative containers: [multi]set, [multi]map

C++11 adds:

0. Addition of `emplace{_front, _back}`.
1. Sequence containers: array, forward_list
2. [nothing new]
3. [nothing new]
4. Unordered associative containers: `unordered_[multi]set`, `unordered_[multi]map`

C++14 adds:

1. Non-member `cbegin/cend/rbegin/rend/cbegin/crend`.



Sequence containers

Ordered associative containers



	STL containers	Some useful operations
Sequence	all containers	size, empty, emplace, erase
	vector<T>	[], at, clear, insert, back, {emplace, push, pop}_back
	deque<T>	[], at, emplace{, _front, _back}, insert, {, push, _pop}_back, {, push, _pop}_front
	list<T>	insert, emplace, merge, reverse, splice, {, emplace, _push, _pop}_back, {back, front}, sort
	array<T>	[], at, front, back, max_size
Associative	forward_list<T>	assign, front, max_size, resize, clear, {insert, erase, emplace}_after, {push, pop, emplace}_front
	set<T>, multiset<T>	find, count, insert, clear, emplace, erase, {lower, upper}_bound
	map<T1, T2>, multimap<T1, T2>	{}, at*, find, count, clear, insert, emplace, erase, {lower, upper}_bound
Unordered Associative	unordered_set<T>, unordered_multiset<T>	find, count, insert, clear, emplace, erase, {lower, upper}_bound
	unordered_map<T1, T2>, unordered_multimap<T1, T2>	{}, at*, find, count, clear, insert, emplace, erase, hash_function
	stack	top, push, pop, swap
Container Adaptors	queue	front, back, push, pop
	priority_queue	top, push, pop, swap
Other	bitset (N bits)	[], count, any, all, none, set, reset, flip

Red means "there's also a stand-alone algorithm of this name"
 Can't iterate over stack, queue, priority_queue.
 * Not on multimap

CS138

1. Sequence Containers

There is a total ordering of contiguous values (i.e., no gaps) on elements based on the *order* in which they are added to the container.

They provide very similar basic functionality, but differ on:

1. Some access methods.

- vector and deque allow random access to elements (via [] / at()), but list allows only sequential access (via iterators).
- deque allows push_back *and* push_front (+ pop_front, + front).

2. Performance.

- vector and deque are optimized for (random access) retrieval, whereas list is optimized for (positional) insertion/deletion.

CS138

U Waterloo CS247 (Spring 2017) — p.14/56

vector<T>

CS138

Can think of as an expandable array that supports access with bounds checking, via `vector<T>::at()`.

Vector elements must be stored contiguously according to the C++ standard, so pointer arithmetic will work and $O(1)$ random access is guaranteed.


- So it pretty much has to be implemented using a C-style array.

Calling `push_back` when vector is at capacity forces a *reallocation*.

Access kind	Complexity	API support
random access	$O(1)$	operator[] or at()
append/delete last	$O(1)^*/O(1)$	push_back/pop_back
prepend/delete first	$O(N)$	not supported as API call
random insert/delete	$O(N)$	insert/erase

(likely) vector Implementation

CS138



deque<T>




== “double-ended queue”; similar to vectors, but allow fast insertion/deletion at beginning and end.

Random access is “fast”, but no guarantee that elements are stored contiguously.

- So pointer arithmetic won’t work.
- `operator[]` and `at()` are overloaded to work correctly.

Access kind	Complexity	API support
random access	$O(1)$	<code>operator[]</code> or <code>at()</code>
append/delete last	$O(1)*O(1)$	<code>push_back/pop_back</code>
prepend/delete first	$O(1)*O(1)$	<code>push_front/pop_front</code>
random insert/delete	$O(N)$	<code>insert/erase</code>

deque implementation



vector vs. deque



So, in real life, should you use `vector` or `deque`?

- If you need to insert at the front, use `deque`.
- If you need to insert in the middle, use `list`.

Random access to elements is constant time in both, but a `vector` may be faster in reality.

Reallocations:

- take longer with a `vector`.
- `vector` invalidates external refs to elements, but not so with a `deque`.
- `vector` copies *elements* (which may be objects), whereas `deque` copies only *ptrs*.

Integrity of External References



```
#include <vector>
#include <deque>
using namespace std;

int main (int argc, char* argv[]) {
    cout << "\nWith a vector:" << endl;
    vector<int> v;
    v.push_back(4);    v.push_back(3);
    v.push_back(37);   v.push_back(15);
    int* p = &v.back();
    cout << *p << " " << v.at(3) << " " // Must be same
         << p << " " << &v.at(3) << endl; // Must be same
    v.push_back(99);    // Causes a reallocation
    cout << *p << " " << v.at(3) << " " // May be different**
         << p << " " << &v.at(3) << endl; // Probably different
```

**** p is no longer pointing to v[3], but the old value of 15 may still "be there"**

Integrity of External References



```
cout << "\nWith a deque:" << endl;
deque<int> d;
d.push_back(4); d.push_back(3);
d.push_back(37); d.push_back(15);
p = &d.back();
cout << *p << " " << d.at(3) << " " // Must be same
    << p << " " << &d.at(3) << endl; // Must be same
d.resize(32767); // Probably causes realloc
cout << *p << " " << d.at(3) << " " // Must be same
    << p << " " << &d.at(3) << endl; // Must be same
}
// My output below, YMMV but comments above will hold
With a vector:
15 15 0x7ff87bc039cc 0x7ff87bc039cc
15 15 0x7ff87bc039cc 0x7ff87bc039ec
With a deque:
15 15 0x7ff87c00220c 0x7ff87c00220c
15 15 0x7ff87c00220c 0x7ff87c00220c
```

list<T>



Implemented as a (plain old) doubly-linked list (PODLL)
Designed for fast insertion and deletion.

Supports only sequential access to elements via iterators.

– No random access via indexing `operator[]` or `at()`.

Access kind	Complexity	API support
random access	O(N)	<i>not supported as an API call</i>
append/delete last	O(1)	<code>push_back/pop_back</code>
prepend/delete first	O(1)	<code>push_front/pop_front</code>
random insert/delete	O(1) (once you've arrived at the elt; O(N) to get there)	<code>insert/erase</code>

(C++11) `std::array`



Effectively, a very thin wrapper around a C++ array, to make it a little more like a fixed-size `vector`.

C++ array vs. `std::array`

- not implicitly converted by compiler into a pointer
- supports many useful functions like
 - an `at()` method, for safe, bounds-checked accessing
 - a `size()` method that returns the extent of the array (which you set when you instantiated the array)

`std::array` vs. `std::vector`

- strong typing: if you know the array size should be fixed, enforce it!
- array contents may be stored on the stack rather than the heap
- `std::array` is faster and more space efficient

(C++11) `std::forward_list`



Basically, a plain-old singly-linked list.

`std::forward_list` vs. `std::list`

- more space efficient, and insertion/deletion operators are slightly faster
- no immediate access to the end of the list
 - no `push_back()`, `back()`
- no ability to iterate backwards
- no `size()` method

2. Container Adapters



Usually a trivial wrapping of a sequence container, to provide a specialized interface with ADT-specific operations to add/remove elements.

- stack, queue, priority_queue

You can specify in the constructor call which container you want to be used in the underlying implementation:

```
stack: vector, deque*, list
queue: deque*, list
priority_queue: vector*, deque
```

[* means default choice]

U Waterloo CS247 (Spring 2017) — p.25/56

```
template <typename T>
class Stack {
public :
    Stack ();
    virtual ~Stack();
    void push (T val);
    void pop ();
    T top ();
    void print ();
    bool isEmpty();
private :
    vector<T> s;
};
```

Note that the STL provides its own definition of Stack: please use that one! This is just an ad hoc example.

```
template <typename T>
Stack<T>::Stack(): s() {}

template <typename T>
Stack<T>::~~Stack () {}

template <typename T>
void Stack<T>::push(T val){
    s.push_back(val);
}

template <typename T>
void Stack<T>::pop () {
    assert (!isEmpty());
    s.pop_back();
}

// ... etc
```



STL Container Adapters



Implemented using the **adapter design pattern**.

1. Define the interface you really want.
e.g., for stack, we want push(), pop()
2. **Instantiate** (don't inherit) a private data-member object from the "workhorse" container class that will do the actual heavy lifting (e.g., vector).
3. Define operations by **delegating** to operations from the workhorse class.

U Waterloo CS247 (Spring 2017) — p.26/56

"The STL Way"

"The STL Way" encourages you to define your own adapter classes, based on the STL container classes, if you have special-purpose needs that are *almost* satisfied by an existing STL class.

- STL doesn't use inheritance or define any methods as virtual.
- Encourages reuse via *adaptation*, rather than inheritance.
- Interface can be exactly what you want, not constrained by inheritance.



U Waterloo CS247 (Spring 2017) — p.28/56

Inheritance vs. Adaptation

Suppose we would like to implement a card game and we want to model a pile of playing cards.

- Actually, a pile of `Card*`, since the cards will be a shared resource and will get passed around.

We want it to support natural `CardPile` capabilities, like `addCard`, `discard`, `merge`, `print`, `etc.`

We also want the client programmer to be able to treat a `CardPile` like a sequential polymorphic container: `iterate`, `find`, `insert`, `erase`.

CS138

U Waterloo CS247 (Spring 2017) — p.29/56

Inheriting from an STL Container

```
// legal, but is it a good idea?
class CardPile : public vector<Card*> {
public:
    // Constructors and destructor
    CardPile ();
    virtual ~CardPile ();
    // Accessors
    void print () const;
    int getHeartsValue () const;
    // Mutators
    void add (Card* card);
    void add (CardPile & otherPile);
    void remove (Card* card);
    void shuffle ();
};
```

CS138

U Waterloo CS247 (Spring 2017) — p.30/56

Traditional STL Adaptation

```
class CardPile {
public:
    // Constructors and destructor
    CardPile();
    virtual ~CardPile();

    // Accessors
    void print() const;
    int getHeartsValue() const;

    // Mutator ops ("natural" for CardPile)
    void add( Card * card );
    void add( CardPile & otherPile );
    void remove( Card * card );
    void shuffle();

    // If want shuffling to be repeatable, pass in a random
    // number generator.
    void shuffle( std::mt19937 & gen );
```

CS138

U Waterloo CS247 (Spring 2017) — p.31/56

Traditional STL adaptation (cont)

```
// Wrapped container methods and types
using iterator = std::vector<Card*>::iterator;
using const_iterator = std::vector<Card*>::const_iterator;
CardPile::iterator begin();
CardPile::const_iterator begin() const;
CardPile::iterator end();
CardPile::const_iterator end() const;
int size() const;
Card * at(int i) const;
void pop_back();
Card * back() const;
bool empty() const;
private:
    std::vector<Card*> pile;
};
// Example of function wrapper
void CardPile::add( CardPile & otherPile ) {
    for ( auto card: otherPile ) pile.emplace_back( card );
    otherPile.pile.clear();
} // CardPile::add
```

CS138

U Waterloo CS247 (Spring 2017) — p.32/56

But there is another way...

Public inheritance:

```
class Circle : public Figure{ ...
```

- Inside the class definition of `Circle`, we have direct access to all non-private members of `Figure`.
- `Circle` is a subtype of `Figure`, and it provides a superset of the `Figure`'s public interface.

Private inheritance:

```
class Circle : private Figure{ ...
```

- Inside the class definition of `Circle`, we have direct access to all non-private members of `Figure`.
- `Circle` is **not** a subtype of `Figure`; it does not support `Figure`'s public interface.
- Client code that instantiates a `Circle` cannot treat it polymorphically as if it were a `Figure`.
 - Cannot invoke any `Figure` public methods.
 - Cannot instantiate a `Circle` to a `Figure*`.

U Waterloo CS247 (Spring 2017) — p.33/56

private Inheritance

Private inheritance is used to allow reuse of a base class's *implementation* without having to support the base class's interface.

All of the inherited `public` (and `protected`) members of the base class are `private` in the child class and can be used to implement child class methods; but they are not exported to the public.

We can selectively make some of the methods of the base class visible to the client code using `using`, as in

- `using Figure::getColour;`
- called **promotion**.

U Waterloo CS247 (Spring 2017) — p.34/56

Private Inheritance of STL Container

```
class CardPile : private std::vector<Card*> {  
  
public:  
    // Constructors and destructor  
    CardPile();  
    virtual ~CardPile();  
  
    // Accessors  
    void print() const;  
    int getHeartsValue() const;  
  
    // Mutator ops ("natural" for CardPile)  
    void add( Card * card );  
    void add( CardPile & otherPile );  
    void remove( Card * card );  
    void shuffle();  
  
    // If want shuffling to be repeatable, pass in a random  
    // number generator.  
    void shuffle( std::mt19937 & gen );
```

U Waterloo CS247 (Spring 2017) — p.35/56

Private Inheritance of STL Container (cont)

```
    // "Promoted" container methods and types  
    using std::vector<Card*>::iterator;  
    using std::vector<Card*>::const_iterator;  
    using std::vector<Card*>::begin;  
    using std::vector<Card*>::end;  
    using std::vector<Card*>::size;  
    using std::vector<Card*>::at;  
    using std::vector<Card*>::pop_back;  
    using std::vector<Card*>::back;  
    using std::vector<Card*>::empty;  
  
};
```

U Waterloo CS247 (Spring 2017) — p.36/56

private Inheritance

This approach is safe *because* it breaks polymorphism!

- Cannot instantiate a `CardPile` to a `vector<Card*>`, so there is no risk of a call to the wrong destructor causing a memory leak.
- The client code cannot accidentally call the wrong version of an inherited non-virtual method.
 - None of the inherited functions are visible to clients unless explicitly made so by using `using` (in which case, the parent definition is used).
 - If you *redefine* an inherited function, the client code will get that version, since they can't see the parent version.

Private inheritance is not conceptually very different from adaptation.

- It requires a little less typing.
- It encourages reuse of the parent class's interface where applicable.

3. Associative Containers

Ordered associative containers

`[multi]map`, `[multi]set`

- The ordering of the elements is based on a *key* value (a piece of the element, e.g., employee records sorted by SIN or name or ...)
 - and not by the order of insertion.
- Implemented using a kind of [binary search tree](#) => lookup is $O(\log N)$.
- Can iterate through container elements "in order".

Unordered associative containers [new in C++11]

`unordered_[multi]map`, `unordered_[multi]set`

- No ordering assumed among the elements.
- Implemented using [hash tables](#) => lookup is $O(1)$.
- Can iterate through container elements, but no particular ordering is assumed.

set<T>

A **set** is a collection of (unique) **values**.

- Typical declaration:
`set<T> s;`
- **T** must support a comparison function with **strict weak ordering**
 - i.e., anti-reflexive, anti-symmetric, transitive.
 - Default is `operator<`
 - Can use a user-defined class, but you must ensure that there is a "reasonable" `operator<` defined or provide an ordering *functor* to the set constructor.

Sets do not allow duplicate elements.

- If you try to insert an element that is already present in the set, the set is unchanged. Return value is a pair `<iterator, bool>`.
 - The *second* of the pair indicates whether the insertion was successful.
 - The *first* of the pair is the position of the new/existing element.

```
// Example with user defined class and operator<
#include <algorithm>
#include <set>
#include <iostream>
#include <string>

using namespace std;

class Student {
public:
    Student (string name, int sNum, double gpa);
    string getName() const;
    int getSNum() const;
    double getGPA() const;
private:
    string name_;
    int sNum_;
    double gpa_;
};

Student::Student(string name, int sNum, double gpa) :
    name_(name), sNum_(sNum), gpa_(gpa) {}
string Student::getName() const {return name_;}
int Student::getSNum() const {return sNum_;}
double Student::getGPA() const {return gpa_;}

bool operator== (const Student& s1, const Student& s2) {
    return (s1.getSNum() == s2.getSNum() &&
            (s1.getName() == s2.getName()) &&
            (s1.getGPA() == s2.getGPA()));
}
```

Equality vs. Equivalence

```
bool operator< (const Student& s1, const Student& s2) {
    return s1.getSNum() < s2.getSNum();
}

ostream& operator<< (ostream &os, const Student& s) {
    os << s.getName() << " " << s.getSNum()
    << " " << s.getGPA();
    return os;
}

int main () {
    Student* pJohn = new Student{ "John Smith", 666, 3.7 };
    Student* pMary = new Student{ "Mary Jones", 345, 3.4 };
    Student* pPeter = new Student{ "Peter Piper", 345, 3.1 }; // Same SNum as Mary

    set<Student> s;
    s.insert(*pJohn);
    s.insert(*pMary);
    s.insert(*pPeter);

    // Will print in numeric order of sNum
    for (auto i=s.begin(); i!=s.end();i++){
        cout << i << endl;
    }
    if ( s.find(*pPeter) != s.end() )
        cout << "Found it with set's find()!" << endl;
    if ( find(s.begin(), s.end(), *pPeter ) != s.end() )
        cout << "Found it with STL algorithm find()" << endl;
}
```

Equivalence

The container search methods (e.g., `find`, `count`, `lower_bound`, ...) will use the following test for equality for elements in ordered associate containers *even if you have your own definition of `operator==`*.

```
if ( !(a<b) && !(b<a) )
```

Equality

Whereas the STL algorithms `find`, `count`, `remove_if` compare elements using `operator==`.

map<key,T>

A **map** maps a **key** to a (unique) **value**.

- Typical declaration:

```
map<T1, T2> m;
```
- T1 is the **key field type**; it must support a comparison function with **strict weak ordering**
 - i.e., anti-reflexive, anti-symmetric, transitive.
 - Default is `operator<`.
 - Can use a user-defined class, but you must ensure that there is a "reasonable" `operator<` defined or provide an ordering functor to the map constructor.
- T2 is the **value field**; can be anything that is copyable and assignable.

Querying Map for Element

Intuitive method (lookup via indexing) will **insert the key if it is not already present**:

```
if ( words [ "bach" ] == 0 )
    // bach not present
```

Alternatively, can use map's `find()` operation to return an **iterator** pointing the queried key/value pair.

```
map<string, int>::iterator it;

it = words.find( "bach" );
if ( it == words.end(...) )
    // bach not present
```

end() is iterator value that points beyond the last element in a collection

```
#include <iostream>
#include <string>
#include <cassert>
#include <map>
using namespace std;

// Examples adapted from Josuttis
int main () {
    map<string,string> dict;

    dict["car"] = "voiture";
    dict["hello"] = "bonjour";
    dict["apple"] = "pomme";

    cout << "Printing simple dictionary" << endl;

    for ( auto i: dict ){
        cout << i.first << ":\t" << i.second << endl;
    }
}
```

How are they Implemented?

[multi]set and [multi]map are usually implemented as a **red-black tree** (see CS240).

- This is a binary search tree that keeps itself reasonably balanced by doing a little bit of work on insert/delete.
- Red-black trees guarantee that lookup/insert/delete are all $O(\log N)$ worst case, which is what the C++ standard requires .
- Optimized search methods (e.g., `find`, `count`, `lower_bound`, `upper_bound`).

Because the containers are automatically sorted, you cannot change the value of an element directly (because doing so might compromise the order of elements).

- There are no operations for direct element access.
- To modify an element, you must remove the old element and insert the new value.

// Examples adapted from Josuttis

```
multimap<string,string> mdict;

mdict.insert(make_pair ("car", "voiture"));
mdict.insert(make_pair ("car", "auto"));
mdict.insert(make_pair ("car", "wagon"));
mdict.insert(make_pair ("hello", "bonjour"));
mdict.insert(make_pair ("apple", "pomme"));

cout << "\nPrinting all defs of \"car\"" << endl;

for (multimap<string,string>::const_iterator
    i=mdict.lower_bound("car");
    i!=mdict.upper_bound("car"); i++) {

    cout << (*i).first << " : " << (*i).second << endl;
}
}
```

(C++11) Unsorted Associative Containers

They are:

```
unordered_[multi]set
unordered_[multi]map
```

They are pretty much the same as the sorted versions except:

- They're not sorted. 😊
- They're implemented using hash tables, so they are $O(1)$ for insert/lookup/remove.
- They do provide iterators that will traverse all of the elements in the container, just not in any "interesting" order.

The C++ Standard Template Library (STL)

The STL is a major component of the C++ Standard Library; it is a large collection of general-purpose generic classes, functions, and iterators:

1. Generic **containers** that take the element type as a parameter.
- e.g., `vector`, `list`, `deque`, `set`, `stack`, `queue`, ...
2. Different kinds of **iterators** that can navigate through the containers.
3. **Algorithms** that (via iterators) perform an interesting operation on a range of elements.
- e.g., `sort`, `random_shuffle`, `transform`, `find`



STL Iterators

The iterator is a fundamental *design pattern*.

- It represents an abstract way of walking through all elements of some interesting data structure.
- You start at the beginning, advance one element at a time, until you reach the end.

In its simplest form, we are given:

- A pointer to the first element in the collection.
- A pointer to just beyond the last element; reaching this element is the stopping criterion for the iteration.
- A way of advancing to the next element (e.g., `operator++`, `operator--`).



STL Containers Provide Iterators

If `c` is a `vector`, `deque`, `list`, `set`, `map`, etc. then

- `c.begin() / c.cbegin()` returns a pointer to the first element
- `c.end() / c.cend()` returns a pointer to just beyond the last element
- `operator++` is defined to advance to the next element

Example

```
vector<string>::const_iterator vi = v.begin();
map<int,string>::iterator mi = mymap.begin();
list<Figure*>::reverse_iterator li = scene.rbegin();
```

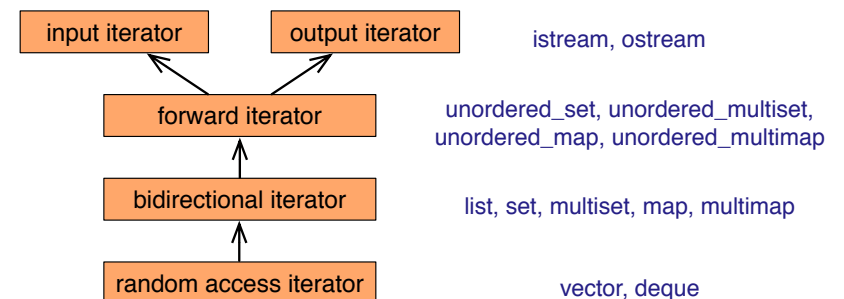
This is the type!

- The iterator types are *nested* types, defined inside the respective container classes, who understand what "++" should mean!



Kinds of Iterators

Iterator **categories** are hierarchical, with lower level categories adding constraints to more general categories.



Why should you care??

Input and Output Iterators

Input iterators are **read-only** iterators where each iterated location may be **read only once**.

Output iterators are **write-only** iterators where each iterated location may be **written only once**.

Operators: `++`, `*` (can be `const`), `==`, `!=` (for comparing iterators)

Mostly used to iterate over streams.

```
#include <iostream>
#include <iterator>
...
copy ( istream_iterator<char> (cin),    // input stream
       istream_iterator<char> (),       // end-of-stream
       ostream_iterator<char> (cout) )  // output stream
```

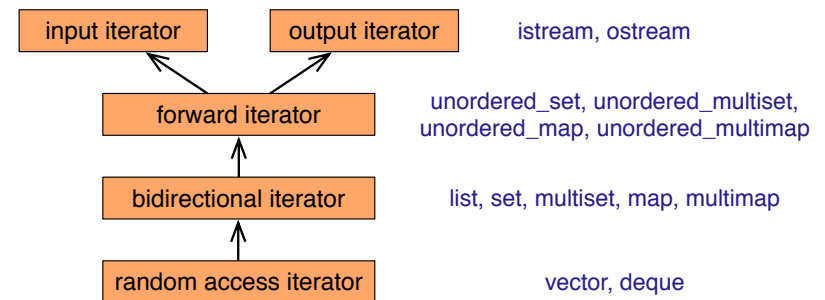
U Waterloo CS247 (Spring 2017) — p.53/56

Other Kinds of Iterators

Forward iterators can read and write to the same location repeatedly.

Bidirectional iterators can iterate backwards (`--`) and forwards (`++`).

Random access iterators : can iterate backwards (`--`) and forwards (`++`), access any element (`[]`), iterator arithmetic (`+`, `-`, `+=`, `-=`).



U Waterloo CS247 (Spring 2017) — p.54/56

Insert Iterators (Inserters)

Iterator that inserts elements into its container:

[back inserter](#): uses container's `push_back()`
[front inserter](#): uses container's `push_front()`
[inserter](#): uses container's `insert()`

```
#include <algorithm>
#include <iterator>
#include <iostream>
#include <vector>
#include <string>

istream_iterator< string > is (cin);
istream_iterator< string > eof;    // end sentinel
vector< string > text;
copy ( is, eof, back_inserter( text ) );
```

U Waterloo CS247 (Spring 2017) — p.55/56

Concluding Remarks

Iterators are a great example of both information hiding and polymorphism.

- Simple, natural, uniform interface for accessing all containers or data structures.
- Can create iterators (STL-derived or homespun) for our own data structures.
- STL iterators are compatible with C pointers, so we can use STL algorithms with legacy C data structures.

U Waterloo CS247 (Spring 2017) — p.56/56